

Chapter 4: Threads & Concurrency

Assignment Project Exam Help

<https://powcoder.com>

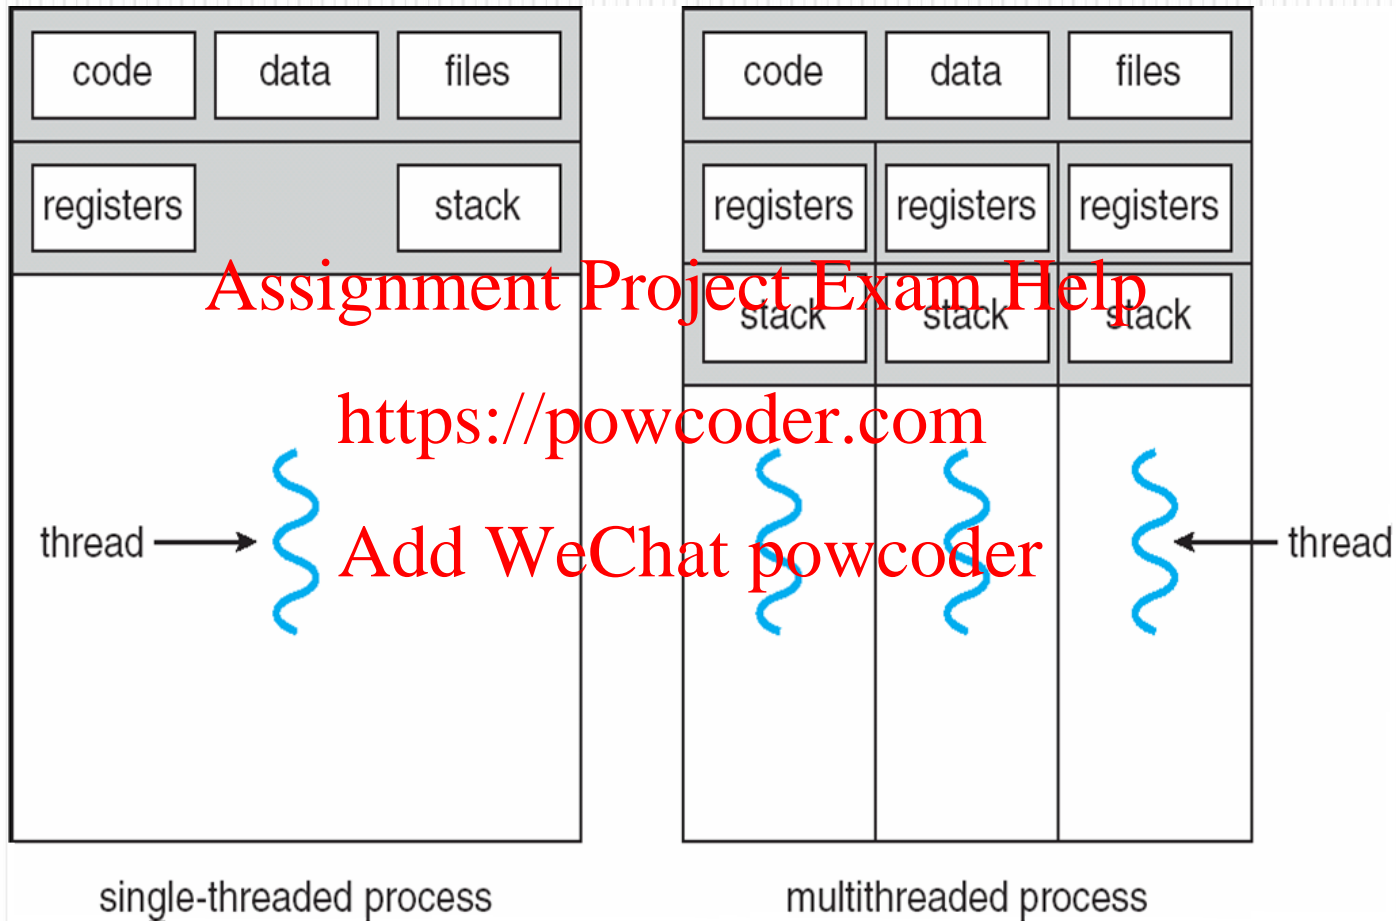
Add WeChat powcoder

Threads

- A thread is a basic unit of CPU utilization; it consists of:
 - its own copy of registers and stack space
 - its own program counter
 - a thread ID
- Threads within a same process or *task* execute in the same address space. A thread shares with its peer threads:
 - code section
 - data section
 - operating system resources, e.g. open files.

A traditional or heavyweight process is equal to a task with one thread.

Single and Multithreaded Processes



User Level Threads

Threads can be implemented at *user* or *kernel* level

- User level threads are implemented in user-level libraries, rather than via system calls, so thread switching does not need to call the operating system and generate an interrupt; kernel has no knowledge of user-level threads

Assignment Project Exam Help

<https://powcoder.com>

User level threads:

Add WeChat powcoder

- Advantage: no kernel involvement in switching; fast
- Disadvantage: since the kernel is not aware of user-level threads, if one user level thread blocks when making a *system call*, all user level threads in the same task may block too

Benefits of Threads

- One multi-threaded process uses fewer resources than multiple redundant processes, including memory, open files and CPU scheduling
 - Low overhead for creation
 - Low overhead for context switch from one thread to another peer thread
- Since threads within the same process share memory, sharing of information between peer threads is easier (but may require the use of synchronization)

Benefits of Threads (Cont.)

- In a multiple threaded task, while one server thread is blocked and waiting, a second thread in the same task can run.
 - Cooperation of multiple threads in same job confers higher throughput and improved performance.
 - Applications that require sharing a common buffer (i.e., producer-consumer) benefit from thread utilization.
- Kernel level threads provide a mechanism that allows sequential processes to make *blocking system calls* while also achieving parallelism.

Potential Pitfalls

Potential pitfalls in using threads

- Because peer threads execute in the same address space, synchronization is often needed to protect shared data against simultaneous access

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

General Benefits of Threads

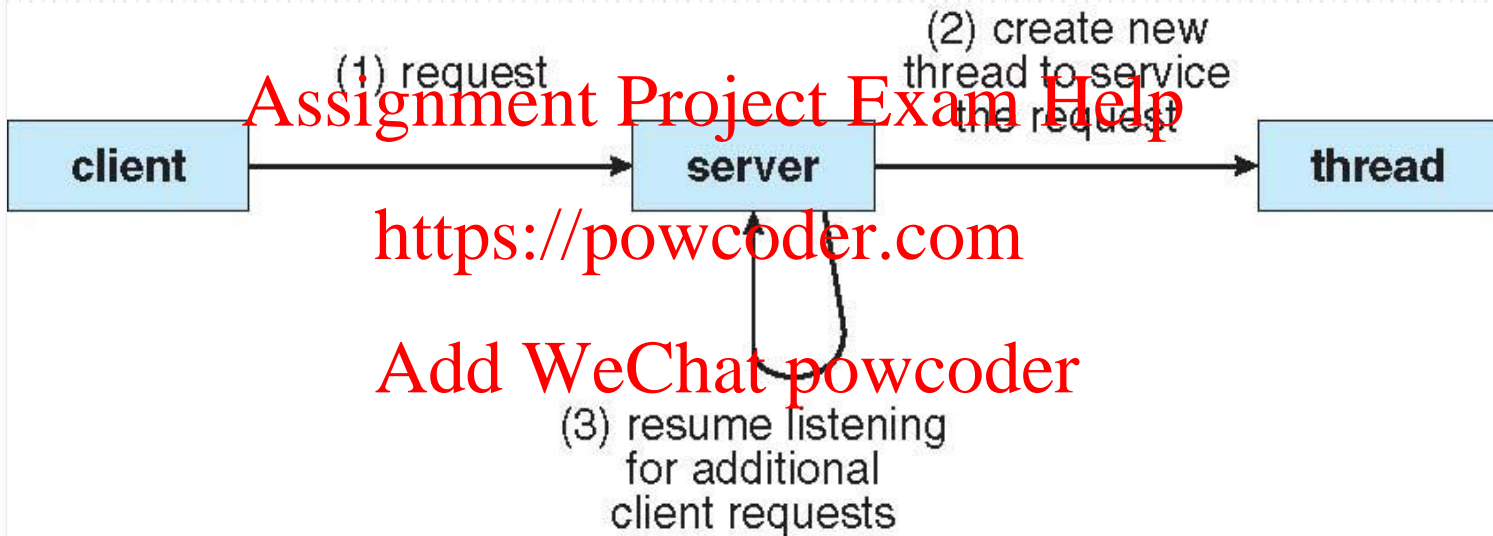
- Responsiveness
 - multithreading an interactive application may allow a program to continue running even if part of it is blocked or is performing a lengthy application.
- Resource Sharing
 - threads share code and data within the same address space.
- Economy
 - it is less costly to create and context switch threads than processes.
- Scalability
 - multithreading on a multi-CPU machine increases parallelism.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


Multithreaded Server Architecture



Multicore Programming Issues

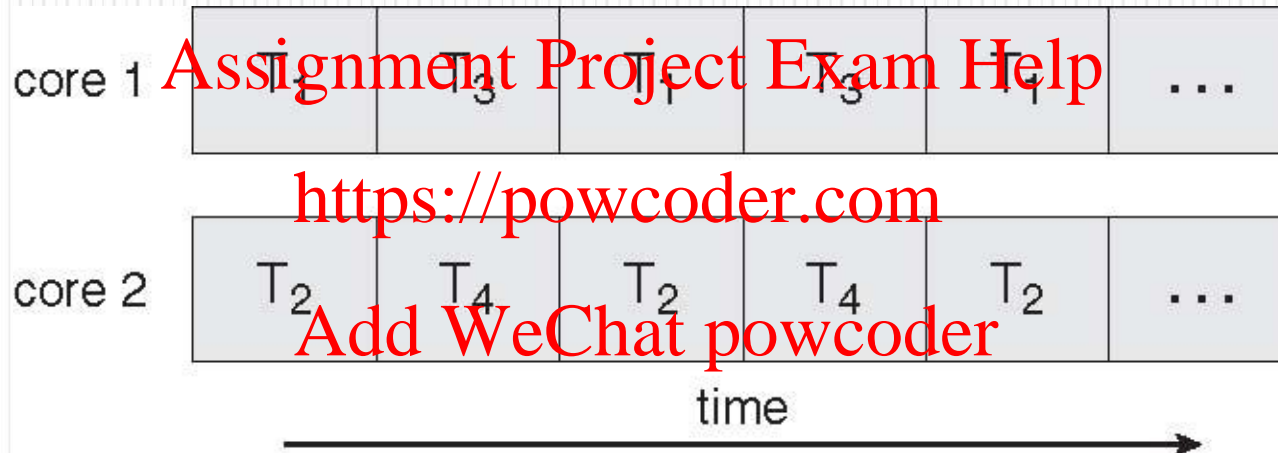
- Dividing activities
 - how to divide an application into separate tasks that can be run in parallel on individual cores.
- Balance
 - need to ensure that the tasks that run in parallel have the same work load.
- Data splitting <https://powcoder.com>
 - need to divide the data that is accessed and manipulated by the tasks that run in parallel.
- Data dependency
 - when one task depends on data produced by another task, the execution of the two tasks must be synchronized.
- Testing and debugging
 - when program runs in parallel, many different possible execution paths.

Concurrent Execution on a Single-core System



Add WeChat powcoder

Parallel Execution on a Multicore System



Multithreading Models

- Many-to-One

- One-to-One

- Many-to-Many

Assignment Project Exam Help

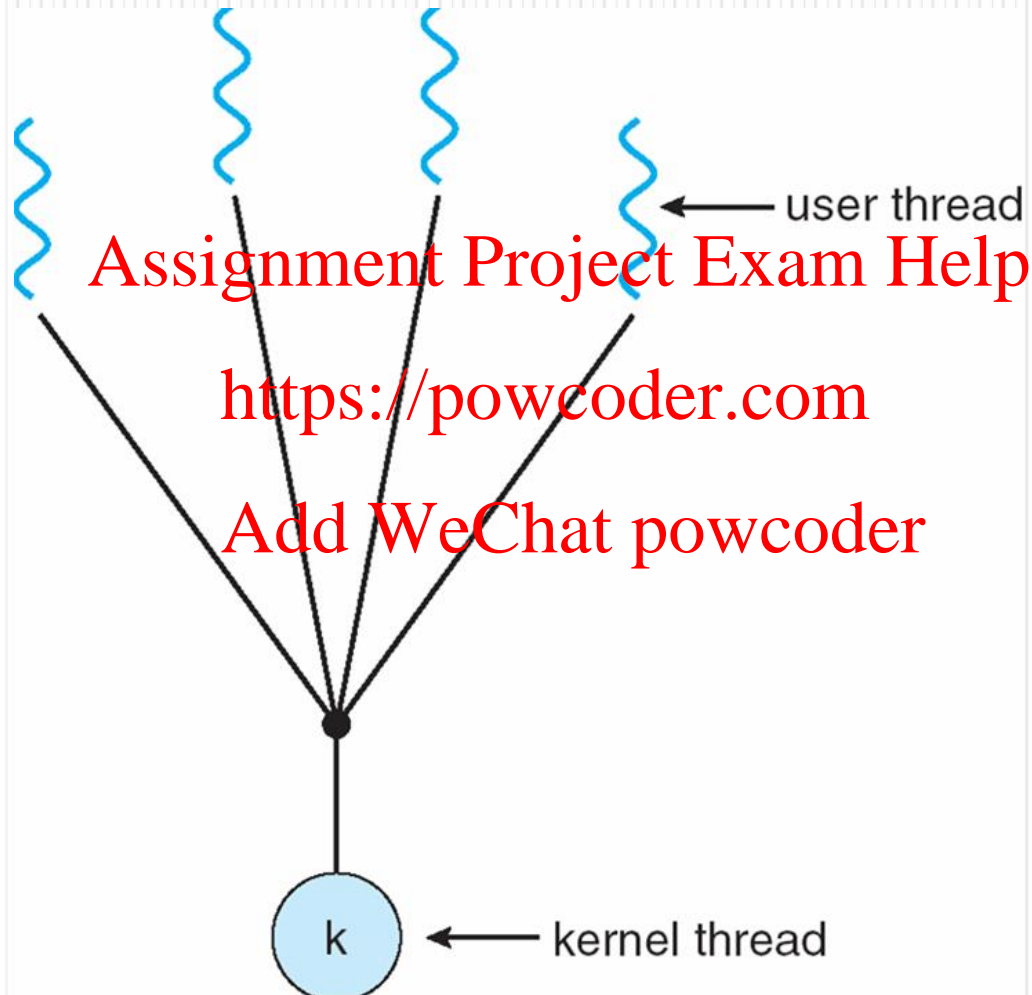
<https://powcoder.com>

Add WeChat powcoder

Many-to-One

- Many user-level threads mapped to single kernel thread
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads
- Advantage: thread management done by thread library in user space, so it is efficient.
- Disadvantages:
 - (1) Entire process will block if one thread makes a blocking system call.
 - (2) Because only one thread can access the kernel at a time, multiple threads are unable to run in parallel on multiprocessors.
- Used on systems that do not support kernel threads.

Many-to-One Model



One-to-One


- Each user-level thread maps to kernel thread
- Examples
 - Windows NT/XP/2000
 - Linux
 - Solaris 9 and later
- Advantages
 - (1) More concurrency – allows another thread to run when a thread makes a blocking system call.
 - (2) Allows multiple threads to run in parallel on multiprocessors.
- Disadvantage:
 - (1) Overhead: creating each user thread requires creating a kernel thread.
 - (2) Need to restrict the number of threads supported by the system.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

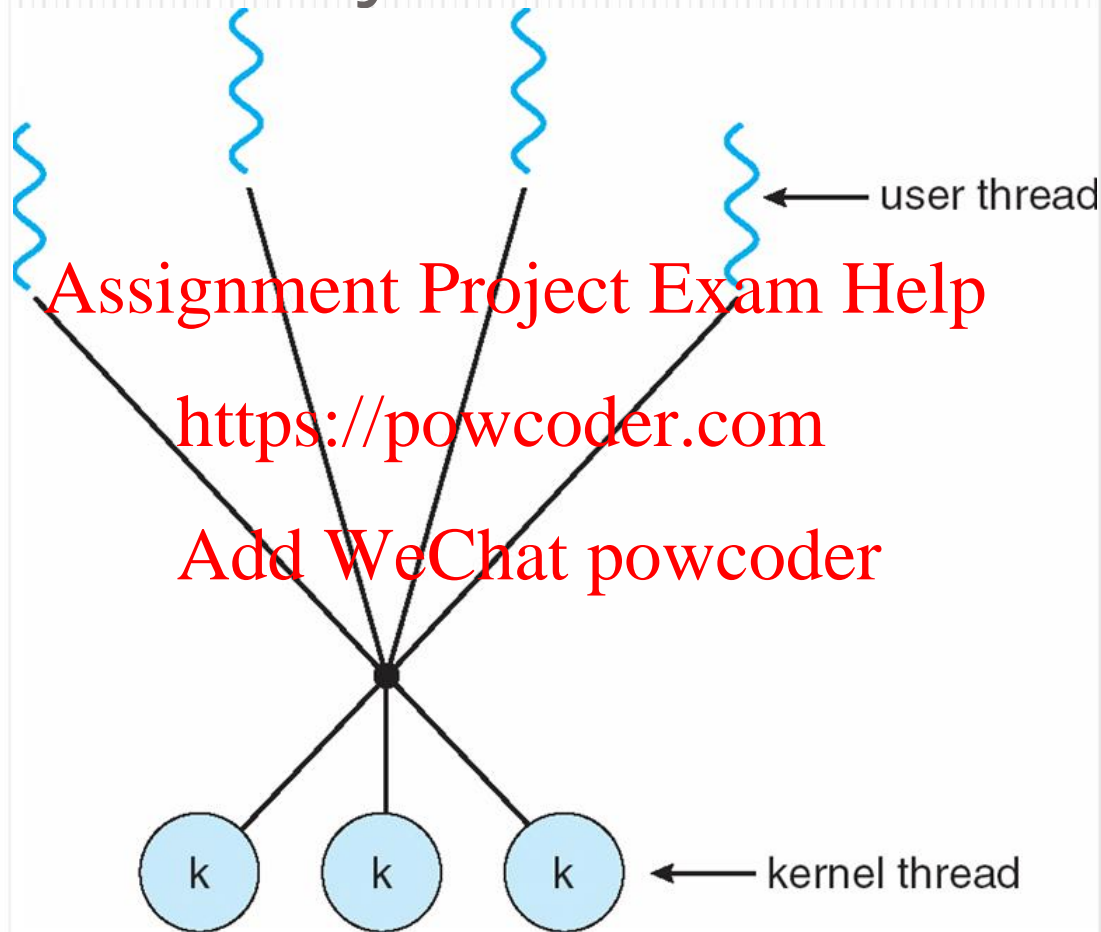
One-to-one Model



Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows NT / 2000 with the Threadlib package
- Compromise between the previous models.
- More complex to implement.

Many-to-Many Model



Thread Libraries

- Thread library provides programmer with API for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Pthreads

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)

Pthread C Program

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_t attr;
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}

/* The thread will execute in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

POSIX pthread Mutexes

- Before accessing shared data, call `pthread_mutex_lock(&mutex)`
 - if another thread has locked the mutex, the caller will be blocked until the other thread unlocks the mutex.
- After accessing shared data, call `pthread_mutex_unlock(&mutex)`
- For nonblocking mutex locks, use `pthread_mutex_trylock(&mutex)`
 - this will return an error status (EBUSY) instead of blocking if the mutex is already blocked
- To prevent simultaneous access to shared data

Linux Threads

- Linux refers to them as *tasks* rather than *threads*
- Thread creation is done through **clone()** system call
<https://powcoder.com>
- **clone()** allows a child task to share the address space of the parent task (process)

End of Chapter 4

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder